

✖ Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from scipy.stats import uniform

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso, Ridge, LogisticRegression
from sklearn.metrics import r2_score, precision_score, recall_score, f1_score
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix, roc_curve
from sklearn.metrics import ConfusionMatrixDisplay, roc_auc_score, precision_recall_curve
from sklearn.preprocessing import OneHotEncoder, LabelEncoder, StandardScaler, MinMaxScaler
from sklearn.impute import KNNImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, cross_validate, GridSearchCV, RandomizedSearchCV
from sklearn.ensemble import GradientBoostingClassifier

from imblearn.over_sampling import SMOTE
from statsmodels.stats.outliers_influence import variance_inflation_factor

#import warnings
#warnings.filterwarnings('ignore')
```

✖ Loading Dataset

```
!wget https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/002/492/original/ola_driver_scaler.csv

--2024-03-19 12:32:57-- https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/002/492/original/ola_driver_scaler.csv
Resolving d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)... 18.172.139.61, 18.172.139.46, 18.172.139.94, ...
Connecting to d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)|18.172.139.61|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1127673 (1.1M) [text/plain]
Saving to: 'ola_driver_scaler.csv'

ola_driver_scaler.c 100%[=====] 1.08M --.-KB/s in 0.06s

2024-03-19 12:32:57 (17.5 MB/s) - 'ola_driver_scaler.csv' saved [1127673/1127673]
```

✖ Analysing the Dataset

```
df = pd.read_csv("ola_driver_scaler.csv")

df.shape

(19104, 14)

df.columns

Index(['Unnamed: 0', 'MMM-YY', 'Driver_ID', 'Age', 'Gender', 'City',
      'Education_Level', 'Income', 'Dateofjoining', 'LastWorkingDate',
      'Joining Designation', 'Grade', 'Total Business Value',
      'Quarterly Rating'],
      dtype='object')

df.head()
```

Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value	
0	0	01/01/19	1	28.0	0.0	C23	2	57387	24/12/18	NaN	1	1	2381061

https://colab.research.google.com/drive/1BvE9EARigs6pNnnAdf1B4X7uNWQAL5DB#scrollTo=kQktKzofhG3r&printMode=true1/19

Next steps:

Generate code with df

1View recommended plots

2	2	03/04/10	1	28.0	0.0	0.23	2	57387	24/12/18	NaN	1	1	-66548
3	3	03/04/10	1	28.0	0.0	0.23	2	57387	24/12/18	03/11/10	1	1	

Basic Data Cleaning and Exploration

Removing Unnecessary Column

```
df.drop(columns = 'Unnamed: 0', axis = 1, inplace = True)
```

Duplicate rows in the dataset

```
df.duplicated().sum()

0
```

Missing values in the dataset

```
df.isnull().sum()

MMM-YY                0
Driver_ID             0
Age                  61
Gender               52
City                 0
Education_Level       0
Income               0
Dateofjoining         0
LastWorkingDate      17488
Joining Designation   0
Grade                0
Total Business Value  0
Quarterly Rating      0
dtype: int64
```

Missing values Treatment

KNN Imputation

Preparing data for KNN imputation



```
#consider only numerical features
num_df = df.select_dtypes('number')

#dropping driver_id(reason is unique)
num_df.drop('Driver_ID', axis = 1, inplace = True)
```

KNN Imputation

```
imputer = KNNImputer(n_neighbors = 5)

num_df = pd.DataFrame(imputer.fit_transform(num_df), columns = num_df.columns)
num_df.head()
```

	Age	Gender	Education_Level	Income	Joining Designation	Grade	Total Business Value	Quarterly Rating	
0	28.0	0.0	2.0	57387.0	1.0	1.0	2381060.0	2.0	
1	28.0	0.0	2.0	57387.0	1.0	1.0	-665480.0	2.0	
2	28.0	0.0	2.0	57387.0	1.0	1.0	0.0	2.0	
3	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0	
4	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0	

Next steps:

[Generate code with num_df](#) [View recommended plots](#)

Concatenating dataframes

```
#remaining columns
rem_col = list(df.columns.difference(num_df.columns))

new_df = pd.concat([df[rem_col], num_df], axis=1)

new_df.shape

(19104, 13)
```

Merging of rows and aggregation of fields

```
data = new_df.groupby(['Driver_ID']).agg({'MMM-YY': 'count', 'Age': 'max', 'Gender': 'max',
                                         'City': 'max', 'Education_Level': 'max', 'Income': 'mean',
                                         'Dateofjoining': 'first', 'LastWorkingDate': 'last', 'Joining Designation': 'max',
                                         'Grade': 'mean', 'Total Business Value': 'sum', 'Quarterly Rating': 'mean'})
```

```
data = data.rename(columns = {'MMM-YY': 'No_of_records', 'Dateofjoining': 'Date_of_joining',
                              'LastWorkingDate': 'Last_working_date', 'Joining Designation': 'Joining_designation',
                              'Total Business Value': 'Total_Business_value', 'Quarterly Rating': 'Quarterly_rating'})
```

```
data['Grade'] = np.round(data['Grade'])
data['Quarterly_rating'] = np.round(data['Quarterly_rating'])
```

```
data.shape

(2381, 12)
```

Column-wise unique entries:

```
for i in data.columns:
    print(f"Unique entries for column {i:20} = {data[i].nunique()}")
```

```
Unique entries for column No_of_records      = 24
Unique entries for column Age                = 61
Unique entries for column Gender             = 6
Unique entries for column City               = 29
Unique entries for column Education_Level    = 3
Unique entries for column Income             = 2339
Unique entries for column Date_of_joining    = 869
Unique entries for column Last_working_date  = 493
Unique entries for column Joining_designation = 5
Unique entries for column Grade              = 5
Unique entries for column Total_Business_value = 1629
Unique entries for column Quarterly_rating   = 4
```

Updating Date-time Columns

```
data['Date_of_joining'] = pd.to_datetime(data['Date_of_joining'])
data['Last_working_date'] = pd.to_datetime(data['Last_working_date'])
```

Conversion of Categorical Attributes to Category

```
categorical_col = ['Gender', 'City', 'Education_Level', 'Joining_designation', 'Grade', 'Quarterly_rating']

for i in categorical_col:
    data[i] = data[i].astype('category')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2381 entries, 1 to 2788
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   No_of_records          2381 non-null   int64
1   Age                    2381 non-null   float64
2   Gender                  2381 non-null   category
3   City                    2381 non-null   category
4   Education_Level         2381 non-null   category
5   Income                  2381 non-null   float64
6   Date_of_joining         2381 non-null   datetime64[ns]
7   Last_working_date       1616 non-null   datetime64[ns]
8   Joining_designation     2381 non-null   category
9   Grade                   2381 non-null   category
10  Total_Business_value    2381 non-null   float64
11  Quarterly_rating        2381 non-null   category
dtypes: category(6), datetime64[ns](2), float64(3), int64(1)
memory usage: 146.4 KB
```

Feature Engineering

```
def check_value(x):
    if len(x) > 1:
        for i in range(len(x)):
            if x[-1] > x[0]:
                return 1
            else:
                return 0
    else:
        return 0
```

Whether the Quarterly Rating has increased for that driver -

for those whose Quarterly Rating has increased we assign the value 1

```
data['Quarterly_rating_increased'] = df.groupby(['Driver_ID'])['Quarterly Rating'].unique().apply(check_value)
```

```
data.Quarterly_rating_increased.value_counts()
```

```
0    1789
1     592
Name: Quarterly_rating_increased, dtype: int64
```

Whether the monthly Income has increased for that driver -

for those whose monthly Income has increased we assign the value 1

```
data['Income_increased'] = df.groupby(['Driver_ID'])['Income'].unique().apply(check_value)
```

```
data.Income_increased.value_counts()
```

```
0    2337
1     44
Name: Income_increased, dtype: int64
```

Target variable creation:

Driver whose last working day is present will have the value 1

```
def check_day(x):
    if x == 0:
        return 0
    else:
        return 1

data['Target'] = (data['Last_working_date'].fillna(0)).apply(check_day)
```

```
data.Target.value_counts()

1    1616
0     765
Name: Target, dtype: int64
```

Number of Months working

Creating column based on their months of working

```
data['No_of_months'] = (data['Last_working_date'].dt.year - data['Date_of_joining'].dt.year) * 12 + (data['Last_working_date'].dt.month - data['Date_of_joining'].dt.month)

data['No_of_months'].fillna(0, inplace = True)
```

Extract Year from Date

```
data['Joining_year'] = data['Date_of_joining'].dt.year
data['Leaving_year'] = data['Last_working_date'].dt.year

data['Joining_year'].fillna(0, inplace = True)
data['Leaving_year'].fillna(0, inplace = True)
```

Final Dataset after removing unnecessary columns

```
data.head()
```

	No_of_records	Age	Gender	City	Education_Level	Income	Date_of_joining	Last_working_date	Joining_designation	Grade	Tr
Driver_ID											
1	3	28.0	0.0	C23	2.0	57387.0	2018-12-24	2019-03-11	1.0	1.0	
2	2	31.0	0.0	C7	2.0	67016.0	2020-11-06	NaT	2.0	2.0	
4	5	43.0	0.0	C13	2.0	65603.0	2019-12-07	2020-04-27	2.0	2.0	
5	3	29.0	0.0	C9	0.0	46368.0	2019-01-09	2019-03-07	1.0	1.0	
6	5	31.0	1.0	C11	1.0	78728.0	2020-07-31	NaT	3.0	3.0	

Next steps:

Generate code with data

☒ View recommended plots



```
f_data = data.drop(columns=['Date_of_joining', 'Last_working_date', 'Quarterly_rating'])
```

```
f_data.shape



(2381, 15)
```

Statistical Summary

```
f_data.describe().T
```

	count	mean	std	min	25%	50%	75%	max	
No_of_records	2381.0	8.023520e+00	6.783590e+00	1.0	3.0	5.0	10.0	24.0	
Age	2381.0	3.377018e+01	5.933265e+00	21.0	30.0	33.0	37.0	58.0	
Income	2381.0	5.923246e+04	2.829821e+04	10747.0	39104.0	55285.0	75835.0	188418.0	
Total_Business_value	2381.0	4.586742e+06	9.127115e+06	-1385530.0	0.0	817680.0	4173650.0	95331060.0	
Quarterly_rating_increased	2381.0	2.486350e-01	4.323126e-01	0.0	0.0	0.0	0.0	1.0	
Income_increased	2381.0	1.847963e-02	1.347062e-01	0.0	0.0	0.0	0.0	1.0	
Target	2381.0	6.787064e-01	4.670713e-01	0.0	0.0	1.0	1.0	1.0	
No_of_months	2381.0	7.941621e+00	1.351620e+01	0.0	0.0	3.0	9.0	85.0	
Joining_year	2381.0	2.018536e+03	1.609597e+00	2013.0	2018.0	2019.0	2020.0	2020.0	
Leaving_year	2381.0	1.370636e+03	9.432429e+02	0.0	0.0	2019.0	2020.0	2020.0	

```
f_data.describe(include = 'category').T
```

	count	unique	top	freq	
Gender	2381.0	6.0	0.0	1382.0	
City	2381	29	C20	152	
Education_Level	2381.0	3.0	2.0	802.0	
Joining_designation	2381.0	5.0	1.0	1026.0	
Grade	2381.0	5.0	2.0	866.0	


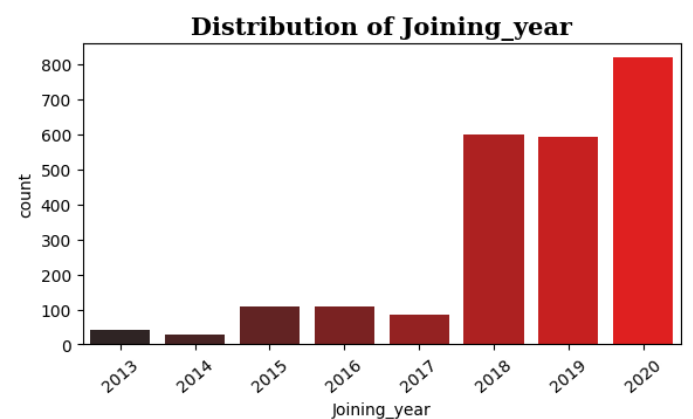



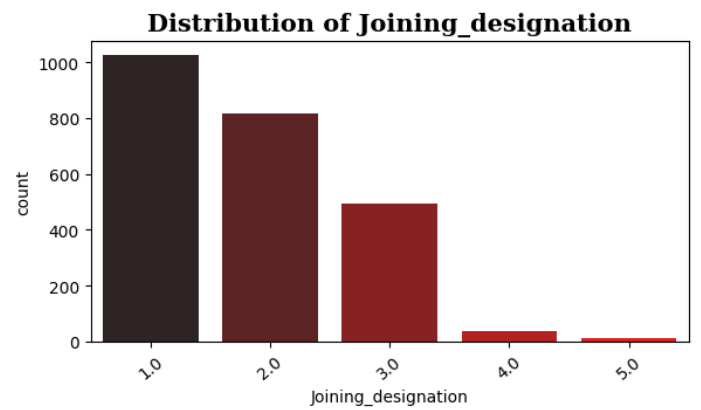
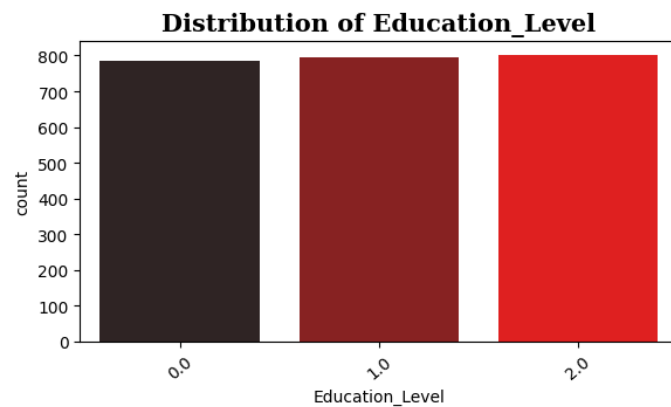


Univariate Analysis

Distribution Plot for Categorical variables

```
categorical = ['Gender', 'City', 'Education_Level', 'Joining_designation',
               'Grade', 'Quarterly_rating_increased', 'Income_increased',
               'Joining_year', 'Leaving_year']

plt.figure(figsize=(12,18))
for i in range(1,10):
    ax=plt.subplot(5,2,i)
    sns.countplot(x=f_data[categorical[i-1]], hue = f_data[categorical[i-1]], legend= False, palette = 'dark:red')
    plt.title(f'Distribution of {categorical[i-1]}', font = 'serif', weight = 'bold', size = 15)
    plt.xticks(rotation = 40)

plt.tight_layout()
plt.show()
```



Distribution Plot For Contineous Variables

```
f_data.select_dtypes('number').columns
```

```
Index(['No_of_records', 'Age', 'Income', 'Total_Business_value',
       'Quarterly_rating_increased', 'Income_increased', 'Target',
       'No_of_months', 'Joining_year', 'Leaving_year'],
      dtype='object')
```

```
contineous = ['Age', 'Income', 'Total_Business_value', 'No_of_months']
```

```
plt.figure(figsize=(12,8))
```

```
for i in range(1,5):
```

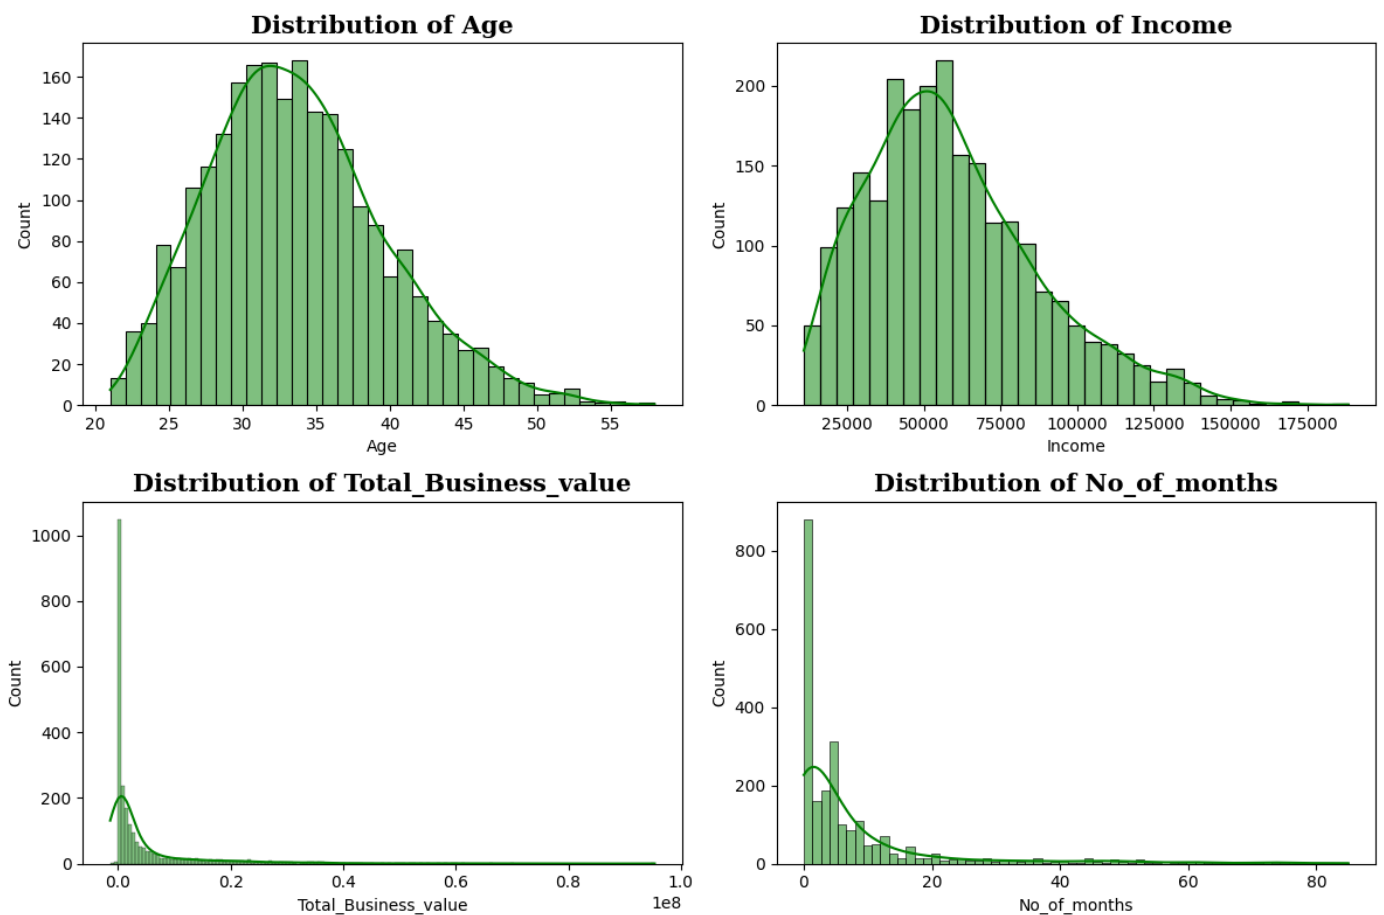
```
    ax=plt.subplot(2,2,i)
```

```
    sns.histplot(f_data[contineous[i-1]], kde = True, color = 'green')
```

```
    plt.title(f'Distribution of {contineous[i-1]}', font = 'serif', weight = 'bold', size = 15)
```

```
plt.tight_layout()
```

```
plt.show()
```

**Bi-variate Analysis****Impact of Categorical variables on Target variable**


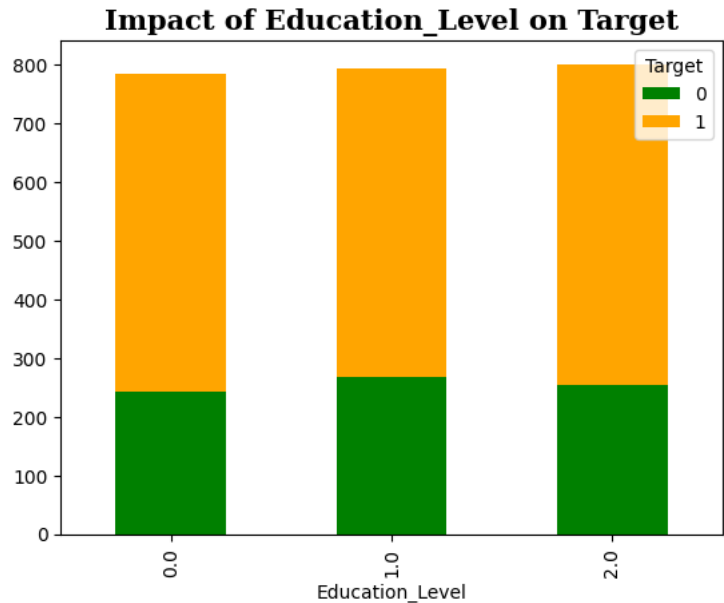

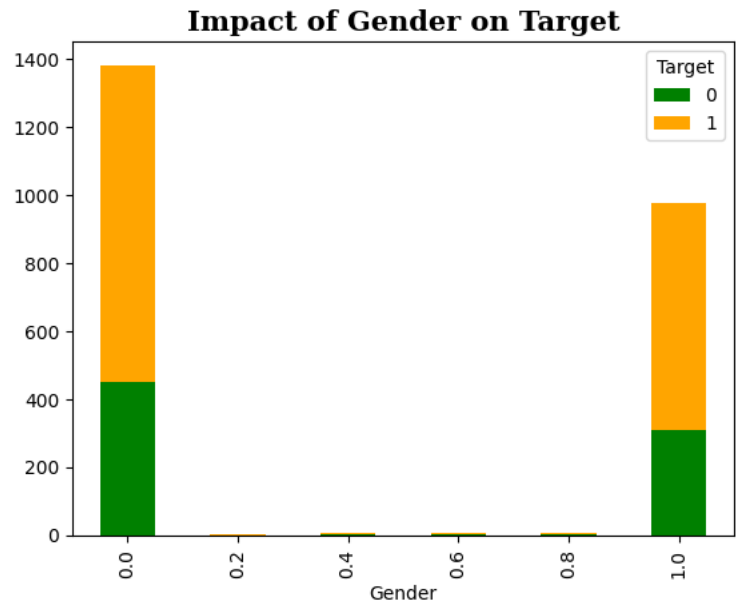
```
#categorical variables
```

```
def bar_plot(cat):
```

```
    bar = f_data.groupby(cat)['Target'].value_counts().unstack().plot(kind='bar', stacked=True, color = ['green', 'orange'])
```



```
cat = ['Gender', 'City', 'Education_Level', 'Joining_designation',  
       'Grade', 'Quarterly_rating_increased', 'Income_increased',  
       'Joining_year']  
for i in cat:  
    bar_plot(i)
```



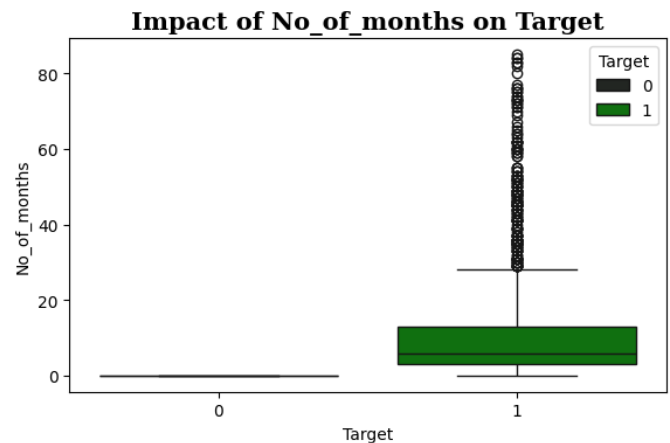
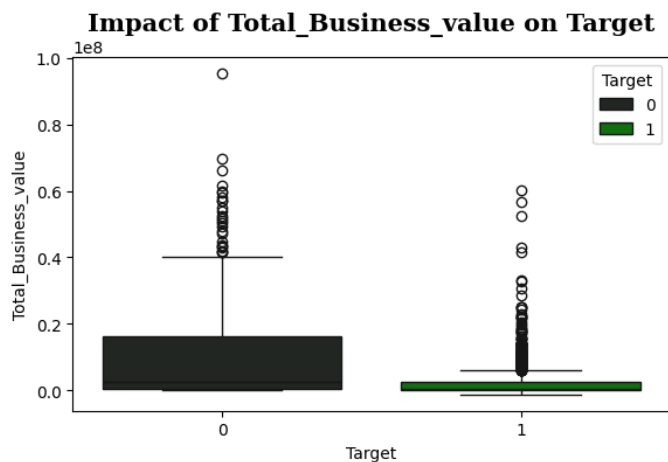
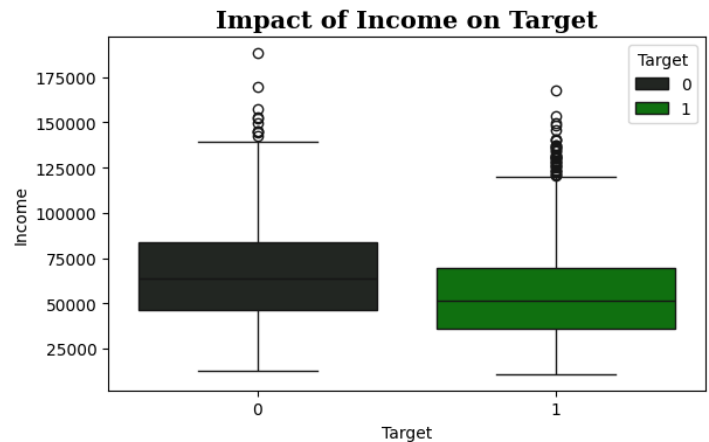
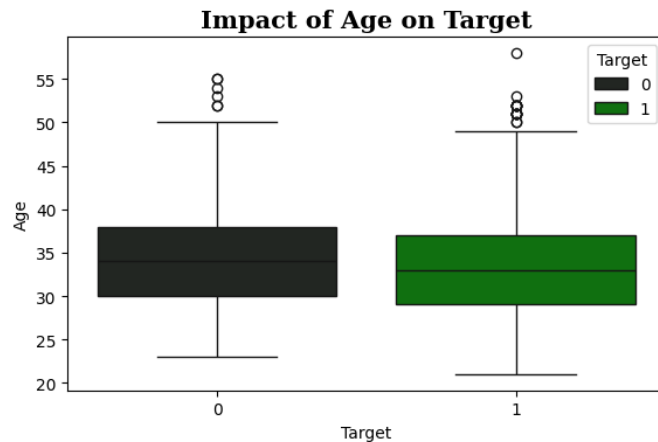
Impact of Contineous variables on Target variable

```
num_col = ['Age', 'Income', 'Total_Business_value', 'No_of_months']

plt.figure(figsize=(12, 8))

for i in range(1, 5):
    ax = plt.subplot(2, 2, i)
    sns.boxplot(x = f_data['Target'], y = f_data[num_col[i-1]], hue = f_data['Target'], palette='dark:green')
    plt.title(f'Impact of {num_col[i-1]} on Target', font = 'serif', weight = 'bold', size = 15)

plt.tight_layout()
plt.show()
```



Multi_variate Analysis

```
plt.figure(figsize = (10, 8))

sns.heatmap(data = f_data.corr(), annot = True, fmt = '.0.2f')

plt.title('Correlation between diffirent variables using heat map', font='serif', size=15, weight='bold')

plt.show()
```

```
<ipython-input-46-847ba243bca9>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version
sns.heatmap(data = f_data.corr(), annot = True, fmt = '.2f')
```

Correlation between different variables using heat map



```
f_data.shape
```

```
(2381, 15)
```

Model Building

Data Preparation for Model Building

```
#pre-processing of Data
x = f_data.drop('Target', axis = 1)
y = f_data['Target']
```

```
print("Shape of x: ", x.shape)
print("Shape of y: ", y.shape)
```

```
Shape of x: (2381, 14)
Shape of y: (2381,)
```

Transforming Categorical Columns by OneHotEncoding

```
x = pd.get_dummies(x, columns = ['City'])
x.head()
```

	No_of_records	Age	Gender	Education_Level	Income	Joining_designation	Grade	Total_Business_value	Quarterly_rating_incr
Driver_ID									
1	3	28.0	0.0	2.0	57387.0	1.0	1.0	1715580.0	
2	2	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	
4	5	43.0	0.0	2.0	65603.0	2.0	2.0	350000.0	
5	3	29.0	0.0	0.0	46368.0	1.0	1.0	120360.0	
6	5	31.0	1.0	1.0	78728.0	3.0	3.0	1265000.0	

5 rows × 42 columns

Split the data

```
#further split the validation-train data to train and test data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state = 1)
```

```
print("Shape of x_train, y_train: ", x_train.shape, y_train.shape)
print("Shape of x_test, y_test: ", x_test.shape, y_test.shape)
```

```
Shape of x_train, y_train: (1904, 42) (1904,)
Shape of x_test, y_test: (477, 42) (477,)
```

Class Imbalance Treatment

Oversampling by SMOTE

```
#checking train data Imbalance
y_train.value_counts()
```

```
1    1287
0     617
Name: Target, dtype: int64
```

```
#oversampling by SMOTE technique
smt = SMOTE()
```

```
#fit SMOTE to training data
X_sm, y_sm = smt.fit_resample(x_train, y_train)
```

```
#After oversampling imbalance check
y_sm.value_counts()
```

```
1    1287
0    1287
Name: Target, dtype: int64
```

Standardization of training data

```
scaler = StandardScaler()

#standardization
X_train = pd.DataFrame(scaler.fit_transform(x_train), columns = x.columns)

X_test = pd.DataFrame(scaler.transform(x_test), columns = x.columns)
```

```
X_train.head()
```

	No_of_records	Age	Gender	Education_Level	Income	Joining_designation	Grade	Total_Business_value	Quarterly_rating_:
0	0.138382	-1.124970	-0.451046	0.006426	-1.104915	-0.972756	-1.153776	-0.380985	
1	-0.596363	-1.291532	-0.857864	-1.217060	-1.482449	-0.972756	-1.153776	-0.516264	
2	-0.302465	-0.958409	1.176224	1.229912	0.150893	0.234628	-0.070012	-0.418969	
3	-0.596363	0.707203	1.176224	-1.217060	-0.227892	-0.972756	-1.153776	-0.497448	
4	2.342618	1.540009	-0.857864	1.229912	-0.623027	-0.972756	-1.153776	0.995936	

5 rows × 42 columns

✓ Ensemble Learning - Bagging Algorithm

Implementation of Random Forest

```
rf_clf = RandomForestClassifier(class_weight='balanced', random_state=7, max_depth=4, n_estimators=100)
```

```
rf_clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(class_weight='balanced', max_depth=4, random_state=7)
```

```
print("Training Score:", rf_clf.score(X_train, y_train)*100)
print("Testing Score:", rf_clf.score(X_test, y_test)*100)
```

```
Training Score: 100.0
Testing Score: 100.0
```

```
y_train_pred = rf_clf.predict(X_train)
y_test_pred = rf_clf.predict(X_test)
```

```
print("f1 Score for Train data:", f1_score(y_train, y_train_pred))
print("f1 Score for Test data:", f1_score(y_test, y_test_pred))
print("Precision Score for Train data:", precision_score(y_train, y_train_pred))
print("Precision Score for Test data:", precision_score(y_test, y_test_pred))
print("Recall Score for Train data:", recall_score(y_train, y_train_pred))
print("Recall Score for Test data:", recall_score(y_test, y_test_pred))
```

```
f1 Score for Train data: 1.0
f1 Score for Test data: 1.0
Precision Score for Train data: 1.0
Precision Score for Test data: 1.0
Recall Score for Train data: 1.0
Recall Score for Test data: 1.0
```

Using Cross validation for better result

```
kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(rf_clf, X_train, y_train, cv=kfold, scoring='accuracy', return_train_score=True)
```

```
print(f"K-Fold Accuracy Mean: \n Train: {cv_acc_results['train_score'].mean()*100:.2f} \n Test: {cv_acc_results['test_score'].mean()*100:.2f} \n")
print(f"K-Fold Accuracy Std: \n Train: {cv_acc_results['train_score'].std()*100:.2f}, \n Test: {cv_acc_results['test_score'].std()*100:.2f} \n")
```

```
K-Fold Accuracy Mean:
Train: 100.00
Test: 100.00
K-Fold Accuracy Std:
Train: 0.00,
Test: 0.00
```

Hyperparameter Tuning

Grid Search CV

```
# Defining parameters -
```

```
params = {
    'n_estimators' : [100,200,300,400],
    'max_depth' : [3,5,10],
    'criterion' : ['gini', 'entropy'],
    'bootstrap' : [True, False],
    'max_features' : [8,9,10]
}
```

```
grid = GridSearchCV(estimator = RandomForestClassifier(),
                    param_grid = params,
                    scoring = 'accuracy',
                    cv = 3,
                    n_jobs=-1
                    )
```

```
grid.fit(X_train, y_train)
```

```
print("Best params: ", grid.best_params_)
print("Best score: ", grid.best_score_)
```

```
Best params: {'bootstrap': True, 'criterion': 'gini', 'max_depth': 3, 'max_features': 8, 'n_estimators': 100}
Best score: 1.0
```

```
rf_clf2 = RandomForestClassifier(class_weight='balanced', random_state=7, bootstrap=True, criterion='gini',
                               max_depth=2, max_features=4, n_estimators=100)
```

```
kfold = KFold(n_splits=10)
```

```
cv_acc_results = cross_validate(rf_clf2, X_train, y_train, cv=kfold, scoring='accuracy', return_train_score=True)
```

```
print(f"K-Fold Accuracy Mean: \n Train: {cv_acc_results['train_score'].mean()*100:.3f} \n Test: {cv_acc_results['test_score'].mean()*100:.3f}")
print(f"K-Fold Accuracy Std: \n Train: {cv_acc_results['train_score'].std()*100:.3f}, \n Test: {cv_acc_results['test_score'].std()*100:.3f}")
```

```
K-Fold Accuracy Mean:
Train: 99.568
Test: 99.633
K-Fold Accuracy Std:
Train: 0.098,
Test: 0.409
```

Randomized Search

```
# Defining parameters -
```

```
params = {'ccp_alpha': uniform(loc=0, scale=0.4)}
```

```
random = RandomizedSearchCV(estimator = RandomForestClassifier(class_weight='balanced', random_state=7, bootstrap=True, criterion='gini',
                                                              max_depth=2, max_features=4, n_estimators=100),
                           param_distributions = params,
                           scoring = 'accuracy',
                           cv = 3,
                           n_iter=15,
                           n_jobs=-1
                           )
```

```
random.fit(X_train, y_train)
```

```
print("Best param: ", random.best_params_)
print("Best score: ", random.best_score_)
```

```
Best param: {'ccp_alpha': 0.2861641672327031}
Best score: 1.0
```

```
rf_clf3 = RandomForestClassifier(class_weight='balanced', random_state=7, bootstrap=True, criterion='gini',
                               max_depth=2, max_features=4, ccp_alpha = 0.28, n_estimators=100)
```

```

kfold = KFold(n_splits=10)
cv_acc_results = cross_validate(rf_clf3, X_train, y_train, cv=kfold, scoring='accuracy', return_train_score=True)

print(f"K-Fold Accuracy Mean: \n Train: {cv_acc_results['train score'].mean()*100:.3f} \n Test: {cv_acc_results['test score'].mean()*100:.3f}")

K-Fold Accuracy Mean:
Train: 99.796
Test: 99.843
K-Fold Accuracy Std:
Train: 0.409,
Test: 0.335

```

Result Evaluation

Classification Report & Confusion Matrix

```
rf_clf3.fit(X_train, y_train)
```

```

RandomForestClassifier
RandomForestClassifier(ccp_alpha=0.04, class_weight='balanced', max_depth=2,
max_features=4, random_state=7)

```

```
y_pred = rf_clf3.predict(X_test)
```

```


print(classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)

```

```
ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=rf_clf3.classes_).plot()
```

	precision	recall	f1-score	support
0	0.99	1.00	1.00	148
1	1.00	1.00	1.00	329
accuracy			1.00	477
macro avg	1.00	1.00	1.00	477
weighted avg	1.00	1.00	1.00	477

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7ec348279ed0>
```



ROC-AUC Curve

```

y_pred = rf_clf3.predict(X_test)
prob = rf_clf3.predict_proba(X_test)
probs = prob[:,1]

```


```

fpr, tpr, thresholds = roc_curve(y_test, probs)
# plot no skill

```




```
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.title("ROC curve")
plt.xlabel('false positive rate')
plt.ylabel('true positive rate')
# show the plot
plt.show()
```



Precision Recall Curve

```
precision, recall, thresholds = precision_recall_curve(y_test, probs)
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
plt.title("Precision Recall curve")
plt.xlabel('Recall')
plt.ylabel('Precision')
# show the plot
plt.show()
```




Comuting Feature Importance

```
# Feature Importance
rf_clf.fit(X_train, y_train)
importances = rf_clf.feature_importances_

indices = np.argsort(importances)[::-1] # Sort feature importances in descending order
names = [X_train.columns[i] for i in indices] # Rearrange feature names so they match the sorted feature importances

plt.figure(figsize=(15, 7)) # Create plot
plt.title("Feature Importance", font = "serif", weight = 'bold', size = 20) # Create plot title
plt.bar(range(X_train.shape[1]), importances[indices]) # Add bars
plt.xticks(range(X_train.shape[1]), names, rotation=90) # Add feature names as x-axis labels
plt.show() # Show plot
```



✓ Ensemble Learning- Boosting Algorithm

Gradient Boosting Classifier

```
parameters = {
    "n_estimators": [50,80,100],
    "max_depth" : [2, 3, 4],
    "max_leaf_nodes" : [5, 10, 20],
    "learning_rate": [0.1, 0.2]
}
```

```
gbc = GradientBoostingClassifier()
```

```
clf = RandomizedSearchCV(gbc, parameters, scoring = "accuracy", cv=3, n_jobs = -1, verbose = 1)
```

```
clf.fit(X_train, y_train)
```

```
clf.fit(X_train, y_train)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```
RandomizedSearchCV
  estimator: GradientBoostingClassifier
    GradientBoostingClassifier
```

```
res = clf.cv_results_

for i in range(len(res["params"])):
    print(f"Parameters:{res['params'][i]} Mean_score: {res['mean_test_score'][i]} Rank: {res['rank_test_score'][i]}")
```

```
Parameters:{'n_estimators': 100, 'max_leaf_nodes': 10, 'max_depth': 4, 'learning_rate': 0.2} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 80, 'max_leaf_nodes': 10, 'max_depth': 4, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 50, 'max_leaf_nodes': 5, 'max_depth': 3, 'learning_rate': 0.2} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 100, 'max_leaf_nodes': 5, 'max_depth': 3, 'learning_rate': 0.2} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 80, 'max_leaf_nodes': 20, 'max_depth': 3, 'learning_rate': 0.2} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 80, 'max_leaf_nodes': 20, 'max_depth': 4, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 50, 'max_leaf_nodes': 10, 'max_depth': 3, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 50, 'max_leaf_nodes': 5, 'max_depth': 4, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 100, 'max_leaf_nodes': 5, 'max_depth': 2, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
Parameters:{'n_estimators': 100, 'max_leaf_nodes': 10, 'max_depth': 4, 'learning_rate': 0.1} Mean_score: 1.0 Rank: 1
```

```
print(clf.best_estimator_)

GradientBoostingClassifier(learning_rate=0.2, max_depth=4, max_leaf_nodes=10)
```

```
gbc = clf.best_estimator_

gbc.fit(X_train, y_train)
```

```
GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.2, max_depth=4, max_leaf_nodes=10)
```

Result Evaluation

Classification Report and Confusion Matrix

```
y_pred = gbc.predict(X_test)
prob = gbc.predict_proba(X_test)
probs = prob[:,1]
```

```
cm = confusion_matrix(y_test, y_pred)
print('Train Score : ', gbc.score(X_train, y_train), '\n')
print('Test Score : ', gbc.score(X_test, y_test), '\n')
print('Accuracy Score : ', accuracy_score(y_test, y_pred), '\n')
print(cm, "---> confusion Matrix ", '\n')
print("ROC-AUC score test dataset: ", roc_auc_score(y_test, prob[:, 1]), '\n')
print("precision score test dataset: ", precision_score(y_test, y_pred), '\n')
print("Recall score test dataset: ", recall_score(y_test, y_pred), '\n')
print("f1 score test dataset: ", f1_score(y_test, y_pred), '\n')
```

Train Score : 1.0

Test Score : 1.0

Accuracy Score : 1.0

```
[[148  0]
 [ 0 329]] ---> confusion Matrix
```

ROC-AUC score test dataset: 1.0